

Вариант 14

Датасет: <https://www.kaggle.com/datasets/noriuk/us-education-datasets-unification-project?resource=download>

набор данных о студентах американских образовательных учреждениях

Задание

Для заданного набора данных (по Вашему варианту) постройте модели классификации или регрессии (в зависимости от конкретной задачи, рассматриваемой в наборе данных). Для построения моделей используйте методы 1 и 2 (по варианту для Вашей группы). Оцените качество моделей на основе подходящих метрик качества (не менее двух метрик). Какие метрики качества Вы использовали и почему? Какие выводы Вы можете сделать о качестве построенных моделей? Для построения моделей необходимо выполнить требуемую предобработку данных: заполнение пропусков, кодирование категориальных признаков, и т.д.

Задание по группам

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing
from sklearn import svm
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import mean_absolute_error,
mean_absolute_percentage_error, mean_squared_error
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, confusion_matrix, classification_report

data = pd.read_csv("drive/MyDrive/Colab Notebooks/states_all.csv",
sep=",")
```

```
data.info()
data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1715 entries, 0 to 1714
```

```
Data columns (total 25 columns):
```

#	Column	Non-Null Count	Dtype
0	PRIMARY_KEY	1715 non-null	object
1	STATE	1715 non-null	object
2	YEAR	1715 non-null	int64
3	ENROLL	1224 non-null	float64
4	TOTAL_REVENUE	1275 non-null	float64
5	FEDERAL_REVENUE	1275 non-null	float64
6	STATE_REVENUE	1275 non-null	float64
7	LOCAL_REVENUE	1275 non-null	float64
8	TOTAL_EXPENDITURE	1275 non-null	float64
9	INSTRUCTION_EXPENDITURE	1275 non-null	float64
10	SUPPORT_SERVICES_EXPENDITURE	1275 non-null	float64
11	OTHER_EXPENDITURE	1224 non-null	float64
12	CAPITAL_OUTLAY_EXPENDITURE	1275 non-null	float64
13	GRADES_PK_G	1542 non-null	float64
14	GRADES_KG_G	1632 non-null	float64
15	GRADES_4_G	1632 non-null	float64
16	GRADES_8_G	1632 non-null	float64
17	GRADES_12_G	1632 non-null	float64
18	GRADES_1_8_G	1020 non-null	float64
19	GRADES_9_12_G	1071 non-null	float64
20	GRADES_ALL_G	1632 non-null	float64
21	AVG_MATH_4_SCORE	565 non-null	float64
22	AVG_MATH_8_SCORE	602 non-null	float64
23	AVG_READING_4_SCORE	650 non-null	float64
24	AVG_READING_8_SCORE	562 non-null	float64

```
dtypes: float64(22), int64(1), object(2)
```

```
memory usage: 335.1+ KB
```

	PRIMARY_KEY	STATE	YEAR	ENROLL	TOTAL_REVENUE
0	1992_ALABAMA	ALABAMA	1992	NaN	2678885.0
1	1992_ALASKA	ALASKA	1992	NaN	1049591.0
2	1992_ARIZONA	ARIZONA	1992	NaN	3258079.0
3	1992_ARKANSAS	ARKANSAS	1992	NaN	1711959.0
4	1992_CALIFORNIA	CALIFORNIA	1992	NaN	26260025.0

STATE_REVENUE LOCAL_REVENUE TOTAL_EXPENDITURE

INSTRUCTION_EXPENDITURE \			
0	1659028.0	715680.0	2653798.0
1481703.0			
1	720711.0	222100.0	972488.0
498362.0			
2	1369815.0	1590376.0	3401580.0
1435908.0			
3	958785.0	574603.0	1743022.0
964323.0			
4	16546514.0	7641041.0	27138832.0
14358922.0			

...	GRADES_4_G	GRADES_8_G	GRADES_12_G	GRADES_1_8_G
GRADES_9_12_G \				
0 ...	57948.0	58025.0	41167.0	NaN
NaN				
1 ...	9748.0	8789.0	6714.0	NaN
NaN				
2 ...	55433.0	49081.0	37410.0	NaN
NaN				
3 ...	34632.0	36011.0	27651.0	NaN
NaN				
4 ...	418418.0	363296.0	270675.0	NaN
NaN				

GRADES_ALL_G	AVG_MATH_4_SCORE	AVG_MATH_8_SCORE
AVG_READING_4_SCORE \		
0	731634.0	208.0
207.0		
1	122487.0	NaN
NaN		
2	673477.0	215.0
209.0		
3	441490.0	210.0
211.0		
4	5254844.0	208.0
202.0		

AVG_READING_8_SCORE	
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

[5 rows x 25 columns]

(data.isnull() | data.empty | data.isna()).sum()

PRIMARY_KEY	0
STATE	0
YEAR	0
ENROLL	491
TOTAL_REVENUE	440
FEDERAL_REVENUE	440
STATE_REVENUE	440
LOCAL_REVENUE	440
TOTAL_EXPENDITURE	440
INSTRUCTION_EXPENDITURE	440
SUPPORT_SERVICES_EXPENDITURE	440
OTHER_EXPENDITURE	491
CAPITAL_OUTLAY_EXPENDITURE	440
GRADES_PK_G	173
GRADES_KG_G	83
GRADES_4_G	83
GRADES_8_G	83
GRADES_12_G	83
GRADES_1_8_G	695
GRADES_9_12_G	644
GRADES_ALL_G	83
AVG_MATH_4_SCORE	1150
AVG_MATH_8_SCORE	1113
AVG_READING_4_SCORE	1065
AVG_READING_8_SCORE	1153
dtype: int64	

##Процент пропусков в каждом признаке

```
target_column = 'AVG_READING_4_SCORE'
data = data[(data[target_column].isna() == False)]
```

```
(data.isnull() | data.empty | data.isna()).sum()
```

PRIMARY_KEY	0
STATE	0
YEAR	0
ENROLL	169
TOTAL_REVENUE	127
FEDERAL_REVENUE	127
STATE_REVENUE	127
LOCAL_REVENUE	127
TOTAL_EXPENDITURE	127
INSTRUCTION_EXPENDITURE	127
SUPPORT_SERVICES_EXPENDITURE	127
OTHER_EXPENDITURE	169
CAPITAL_OUTLAY_EXPENDITURE	127
GRADES_PK_G	81
GRADES_KG_G	76
GRADES_4_G	76

GRADES_8_G	76
GRADES_12_G	76
GRADES_1_8_G	209
GRADES_9_12_G	158
GRADES_ALL_G	76
AVG_MATH_4_SCORE	129
AVG_MATH_8_SCORE	129
AVG_READING_4_SCORE	0
AVG_READING_8_SCORE	88

dtype: int64

```
row_count = data.shape[0]
for key, elem in (data.isnull() | data.empty |
data.isna()).sum().items():
    print(key, "{:.2f}".format(elem / row_count * 100) + "%")
```

```
PRIMARY_KEY 0.00%
STATE 0.00%
YEAR 0.00%
ENROLL 26.00%
TOTAL_REVENUE 19.54%
FEDERAL_REVENUE 19.54%
STATE_REVENUE 19.54%
LOCAL_REVENUE 19.54%
TOTAL_EXPENDITURE 19.54%
INSTRUCTION_EXPENDITURE 19.54%
SUPPORT_SERVICES_EXPENDITURE 19.54%
OTHER_EXPENDITURE 26.00%
CAPITAL_OUTLAY_EXPENDITURE 19.54%
GRADES_PK_G 12.46%
GRADES_KG_G 11.69%
GRADES_4_G 11.69%
GRADES_8_G 11.69%
GRADES_12_G 11.69%
GRADES_1_8_G 32.15%
GRADES_9_12_G 24.31%
GRADES_ALL_G 11.69%
AVG_MATH_4_SCORE 19.85%
AVG_MATH_8_SCORE 19.85%
AVG_READING_4_SCORE 0.00%
AVG_READING_8_SCORE 13.54%
```

###Заполнять пропуски в признаках, с процентом пустых значений более 30 имеет мало смысла, поэтому составим список признаков, которые нужно удалить.

```
row_count = data.shape[0]
columns_to_remove = []
for key, elem in (data.isnull() | data.empty |
data.isna()).sum().items():
    if elem / row_count >= 0.3:
```

```

    columns_to_remove.append(key)
columns_to_remove

['GRADES_1_8_G']

data_new = data.drop(columns_to_remove, axis=1)
data_new.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 650 entries, 0 to 1714
Data columns (total 24 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   PRIMARY_KEY                             650 non-null    object
1   STATE                                    650 non-null    object
2   YEAR                                    650 non-null    int64
3   ENROLL                                  481 non-null    float64
4   TOTAL_REVENUE                          523 non-null    float64
5   FEDERAL_REVENUE                        523 non-null    float64
6   STATE_REVENUE                          523 non-null    float64
7   LOCAL_REVENUE                          523 non-null    float64
8   TOTAL_EXPENDITURE                      523 non-null    float64
9   INSTRUCTION_EXPENDITURE                523 non-null    float64
10  SUPPORT_SERVICES_EXPENDITURE            523 non-null    float64
11  OTHER_EXPENDITURE                       481 non-null    float64
12  CAPITAL_OUTLAY_EXPENDITURE              523 non-null    float64
13  GRADES_PK_G                             569 non-null    float64
14  GRADES_KG_G                             574 non-null    float64
15  GRADES_4_G                              574 non-null    float64
16  GRADES_8_G                              574 non-null    float64
17  GRADES_12_G                             574 non-null    float64
18  GRADES_9_12_G                           492 non-null    float64
19  GRADES_ALL_G                            574 non-null    float64
20  AVG_MATH_4_SCORE                        521 non-null    float64
21  AVG_MATH_8_SCORE                        521 non-null    float64
22  AVG_READING_4_SCORE                     650 non-null    float64
23  AVG_READING_8_SCORE                     562 non-null    float64
dtypes: float64(21), int64(1), object(2)
memory usage: 127.0+ KB

```

Обработка пропусков данных

Посмотрим средние значения в каждом столбце с пропусками

```

for column in data_new:
    if (data_new[column].isnull() | data_new[column].empty |
        data_new[column].isna()).sum() != 0:
        if data_new[column].dtype != 'object':
            print(column + ":", data_new[column].describe()[['mean']][0])
        else:
            print(column + ":", data_new[column].describe()[['top']][0])

```

ENROLL: 934450.4885654886
 TOTAL_REVENUE: 9610588.969407266
 FEDERAL_REVENUE: 840781.4225621414
 STATE_REVENUE: 4467784.89292543
 LOCAL_REVENUE: 4302022.653919694
 TOTAL_EXPENDITURE: 9730734.829827916
 INSTRUCTION_EXPENDITURE: 5039264.764818355
 SUPPORT_SERVICES_EXPENDITURE: 2834123.634799235
 OTHER_EXPENDITURE: 461512.7006237006
 CAPITAL_OUTLAY_EXPENDITURE: 954404.3078393881
 GRADES_PK_G: 21156.2460456942
 GRADES_KG_G: 71505.45296167248
 GRADES_4_G: 72367.06794425087
 GRADES_8_G: 72312.85365853658
 GRADES_12_G: 62889.35017421603
 GRADES_9_12_G: 286875.8699186992
 GRADES_ALL_G: 961249.6393728222
 AVG_MATH_4_SCORE: 237.39539347408828
 AVG_MATH_8_SCORE: 279.9846449136276
 AVG_READING_8_SCORE: 263.55871886120997

Заполним пропуски

```

for column in data_new:
    if (data_new[column].isnull() | data_new[column].empty |
        data_new[column].isna()).sum() != 0:
        if data_new[column].dtype != 'object' and column != target_column:
            data_new[column].fillna(data_new[column].median(), inplace=True)

(data_new.isnull() | data_new.empty | data_new.isna()).sum()
  
```

PRIMARY_KEY	0
STATE	0
YEAR	0
ENROLL	0
TOTAL_REVENUE	0
FEDERAL_REVENUE	0
STATE_REVENUE	0
LOCAL_REVENUE	0
TOTAL_EXPENDITURE	0
INSTRUCTION_EXPENDITURE	0
SUPPORT_SERVICES_EXPENDITURE	0
OTHER_EXPENDITURE	0
CAPITAL_OUTLAY_EXPENDITURE	0
GRADES_PK_G	0
GRADES_KG_G	0
GRADES_4_G	0
GRADES_8_G	0
GRADES_12_G	0
GRADES_9_12_G	0
GRADES_ALL_G	0

```

AVG_MATH_4_SCORE          0
AVG_MATH_8_SCORE          0
AVG_READING_4_SCORE       0
AVG_READING_8_SCORE       0
dtype: int64

```

```
data_new.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 650 entries, 0 to 1714
```

```
Data columns (total 24 columns):
```

#	Column	Non-Null Count	Dtype
0	PRIMARY_KEY	650 non-null	object
1	STATE	650 non-null	object
2	YEAR	650 non-null	int64
3	ENROLL	650 non-null	float64
4	TOTAL_REVENUE	650 non-null	float64
5	FEDERAL_REVENUE	650 non-null	float64
6	STATE_REVENUE	650 non-null	float64
7	LOCAL_REVENUE	650 non-null	float64
8	TOTAL_EXPENDITURE	650 non-null	float64
9	INSTRUCTION_EXPENDITURE	650 non-null	float64
10	SUPPORT_SERVICES_EXPENDITURE	650 non-null	float64
11	OTHER_EXPENDITURE	650 non-null	float64
12	CAPITAL_OUTLAY_EXPENDITURE	650 non-null	float64
13	GRADES_PK_G	650 non-null	float64
14	GRADES_KG_G	650 non-null	float64
15	GRADES_4_G	650 non-null	float64
16	GRADES_8_G	650 non-null	float64
17	GRADES_12_G	650 non-null	float64
18	GRADES_9_12_G	650 non-null	float64
19	GRADES_ALL_G	650 non-null	float64
20	AVG_MATH_4_SCORE	650 non-null	float64
21	AVG_MATH_8_SCORE	650 non-null	float64
22	AVG_READING_4_SCORE	650 non-null	float64
23	AVG_READING_8_SCORE	650 non-null	float64

```
dtypes: float64(21), int64(1), object(2)
```

```
memory usage: 127.0+ KB
```

```
category_cols = ['PRIMARY_KEY', 'STATE']
```

```
print("Количество уникальных значений\n")
```

```
for col in category_cols:
```

```
    print(f'{col}: {data_new[col].unique().size}')
```

```
Количество уникальных значений
```

```
PRIMARY_KEY: 650
```

```
STATE: 53
```

```
data_new = data_new.drop(category_cols, axis=1)
```



```

X_data = data_new[['YEAR',
'ENROLL',
'FEDERAL_REVENUE',
'STATE_REVENUE',
'LOCAL_REVENUE',
'INSTRUCTION_EXPENDITURE',
'SUPPORT_SERVICES_EXPENDITURE',
'OTHER_EXPENDITURE',
'CAPITAL_OUTLAY_EXPENDITURE',
'GRADES_ALL_G']]
Y_data = data_new[target_column].to_list()

X_data['YEAR'] = data_new['YEAR'] - data_new['YEAR'].mean()

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#
returning-a-view-versus-a-copy
    """Entry point for launching an IPython kernel.

X_data = preprocessing.normalize(X_data,axis = 0)

X_data.shape

(650, 10)

X_train, X_test, y_train, y_test = train_test_split(
    X_data, Y_data, test_size=0.3, random_state=42)

data_new.isnull().sum()

YEAR                                0
ENROLL                              0
TOTAL_REVENUE                       0
FEDERAL_REVENUE                     0
STATE_REVENUE                       0
LOCAL_REVENUE                       0
TOTAL_EXPENDITURE                   0
INSTRUCTION_EXPENDITURE              0
SUPPORT_SERVICES_EXPENDITURE         0
OTHER_EXPENDITURE                   0
CAPITAL_OUTLAY_EXPENDITURE           0
GRADES_PK_G                         0
GRADES_KG_G                         0
GRADES_4_G                          0
GRADES_8_G                          0
GRADES_12_G                         0
GRADES_9_12_G                       0

```

```

GRADES_ALL_G          0
AVG_MATH_4_SCORE      0
AVG_MATH_8_SCORE      0
AVG_READING_4_SCORE   0
AVG_READING_8_SCORE   0
dtype: int64

```

SVC

Выбрал MAE, MSE и MAPE для оценки качества, так как это самые показательные метрики для регрессии

```

svm_model = SVC()
mape = -
cross_val_score(svm_model,X_train,y_train,cv=4,scoring='neg_mean_absolute_percentage_error').mean()
mae = -
cross_val_score(svm_model,X_train,y_train,cv=4,scoring='neg_mean_absolute_error').mean()
mse = -
cross_val_score(svm_model,X_train,y_train,cv=4,scoring='neg_mean_squared_error').mean()
print('SVM Errors')
print('MAE:' + str(round(mae,3)) + ' MAPE:' + str(round(mape,3)) + ' MSE:' + str(round(mse,3)))

```

The least populated class in y has only 1 members, which is less than n_splits=4.

The least populated class in y has only 1 members, which is less than n_splits=4.

The least populated class in y has only 1 members, which is less than n_splits=4.

SVM Errors

MAE:5.89 MAPE:0.028 MSE:65.313

```

svm_model.fit(X_train,y_train)
mae = mean_absolute_error(y_test,svm_model.predict(X_test))
mape =
mean_absolute_percentage_error(y_test,svm_model.predict(X_test))
mse = mean_squared_error(y_test,svm_model.predict(X_test))
print('MAE:' + str(round(mae,3)) + ' MAPE:' + str(round(mape,3)) + ' MSE:' + str(round(mse,3)))

```

MAE:5.462 MAPE:0.026 MSE:63.851

RandomForestClassifier

```

rfc_model = RandomForestClassifier()
mape = -
cross_val_score(rfc_model,X_train,y_train,cv=2,scoring='neg_mean_absolute_percentage_error').mean()

```

```

ute_percentage_error').mean()
mae = -
cross_val_score(rfc_model,X_train,y_train,cv=2,scoring='neg_mean_absolute_error').mean()
mse = -
cross_val_score(rfc_model,X_train,y_train,cv=2,scoring='neg_mean_squared_error').mean()
print('SVM Errors')
print('MAE:' + str(round(mae,3)) + ' MAPE:' + str(round(mape,3)) + '
MSE:' + str(round(mse,3)))

```

The least populated class in y has only 1 members, which is less than n_splits=2.

The least populated class in y has only 1 members, which is less than n_splits=2.

The least populated class in y has only 1 members, which is less than n_splits=2.

SVM Errors

MAE:4.264 MAPE:0.019 MSE:37.903

```

rfc_model.fit(X_train,y_train)
mae = mean_absolute_error(y_test,rfc_model.predict(X_test))
mape =
mean_absolute_percentage_error(y_test,rfc_model.predict(X_test))
mse = mean_squared_error(y_test,rfc_model.predict(X_test))
print('MAE:' + str(round(mae,3)) + ' MAPE:' + str(round(mape,3)) + '
MSE:' + str(round(mse,3)))

```

MAE:3.851 MAPE:0.018 MSE:34.333

Можно сделать вывод о том, что обе модели очень хорошо показали себя на тестовых и тренировочных данных. Даже при условии что значений было немного.